

Week 2 Exercise: Local diversity

Z620: Quantitative Biodiversity, Indiana University

November 8, 2014

Overview

In this exercise, we will use R and RStudio to explore aspects of local or site-specific diversity, also known as alpha (α) diversity. We will use the tropical forest dataset from Barro-Colorado Island (BCI), a 1,560-hectare island in the center of the Panama Canal that is administered by the Smithsonian. After setting up our working environment, we will explore the BCI dataset and ways to quantify and examine the α -diversity within individual plots of tropical trees.

RETRIEVE AND SET YOUR WORKING DIRECTORY

```
## [1] "/Users/lisalocey/GitHub/Quantitative_Biodiversity/Assignments/Week2"
```

INSTALL PACKAGES

Ecologists have developed many packages for conducting quantitative analyses in R. You can use the ‘help’ function to learn more about package installation and add-ons. In this class, we will often use the package called ‘vegan’, which contains tools commonly used in ecological research, including analyses of (α)-diversity. To install vegan into your R environment type: `install.packages("vegan")` into your RStudio console.

```
require("vegan") #You can also type: library("vegan"). The difference being that "require"
```

```
## Loading required package: vegan
## Loading required package: permute
## Loading required package: lattice
## This is vegan 2.0-10
```

```
#loads additional packages (dependencies) that vegan requires for some analyses.
```

LOADING BCI FROM VEGAN

Conveniently, the vegan library contains a version of the BCI dataset, which is obtained by censusing the trees of 50 1-hectare plots every several years. More information on the BCI forest dataset can be found here: <http://ctfs.arnarb.harvard.edu/webatlas/datasets/bci/>.

```
data(BCI) # Loading the BCI dataset
```

EXPLORING THE SITE-BY-SPECIES MATRIX

In vegan, the BCI data are organized into a **site-by-species matrix**, that is, a data table where the abundances of all 225 tree species (columns) are given for each of the 50 sites (rows). Consequently, this table should have 50 rows (sites) and 225 columns (species). Let’s verify this using the dimension or “dim” function:

```
dim(BCI) # Prints the dimensions of the site (row) by species (column) matrix
```

```
## [1] 50 225
```

The **site-by-species matrix** is one of the most basic data structures used by ecologists. It contains the abundances (or presence and absence) of all species found among a given number of sites, quadrats, transects, etc. With the BCI site-by-species matrix loaded, let's print the abundances of 4 species found in the first 8 sites.

```
BCI[1:8, 14:17] # print abundances for sites (rows) 1 to 8 and for species (columns) 14 to 17
```

```
## Andira.inermis Annona.spraguei Apeiba.aspera Apeiba.tibourbou
## 1              0              1              13              2
## 2              0              0              12              0
## 3              0              1              6              1
## 4              0              0              3              1
## 5              1              0              4              0
## 6              1              0              10             0
## 7              0              0              5              0
## 8              0              1              4              1
```

Here, we can see that Cabbage Bark (*Andira inermis*) is absent from six of the eight sites and is only found as a single individual in two sites. On the other hand, *Apeiba aspera* (locally known as Monkey Comb) is found at all eight sites and is relatively much more abundant than any of the other three species that we selected.

EXPLORING LOCAL DIVERSITY

So far, you have been introduced to an important data structure, the site-by-species matrix. You have learned to print its dimensions and to print the abundances of particular species (columns) found in specific sites (rows). Now, let's focus on a specific site by selecting the first row and assigning it to a variable called "Site1":

```
Site1 <- BCI[1,] # assign the first row (site) to the variable Site1
dim(Site1) # print the dimensions of Site1 to the screen
```

```
## [1] 1 225
```

You can see that Site1 has one row (it is only 1 site) and 225 columns (225 potentially present species). Remember that the BCI site-by-species matrix also accounts for absences, i.e., with zeros. Let's find out how many species actually occupied Site1 when it was censused:

```
S1 <- specnumber(Site1) # Find the number of species in Site1 and assign it to a variable S1
cat('There are', S1, 'species at Site1') # "cat()" concatenates strings and numbers # Mario, feel free
```

```
## There are 93 species at Site1
```

Now we know that the recorded species richness (usually denoted as S) of Site1 is 93. Local species richness is the most basic aspect of diversity and is, in fact, what we usually refer to as α -diversity. Beyond knowing the number of species, we are often interested in knowing how many individuals were found among them. That is, the number of individuals recorded at Site1.

```
N1 <- sum(Site1) # find the number of individuals at Site1 and assign it to a variable N1
cat('There are', N1, 'individuals among', S1, 'species at Site1') # Mario, feel free to edit this as you like

## There are 448 individuals among 93 species at Site1
```

Now we know the recorded species richness (S) and recorded total abundance (usually denoted as N) of Site1 in the BCI site-by-species matrix. N and S are two of the most common pieces of data we obtain when sampling ecological communities. Finally, we are beginning to quantify our data!

The species abundance distribution


A next natural step in exploring aspects of local diversity is simply to ask how abundance varies among species. One of the most common ways to visually explore this relationship, is simply to rank the species from greatest to least, that is, as a rank-abundance curve (RAC). Let's plot the RAC for Site1.

Let's begin by first removing the zeros from Site1. This removes the species that, while in the BCI site-by-species matrix, were not found at Site1.

```
Site1 <- Site1[ !Site1 %in% c(0) ] # Remove all zeros from Site1 and reassign the new vector
```

Now, we can plot the RAC, accounting only for the species that were recorded at Site1.

```
RAC <- Site1[order(Site1, decreasing=TRUE)] # Rank the taxa by abundance (y-values)
ranks <- seq(1, ncol(RAC)) # Create a sequence of S ranks (x-values)
plot(ranks, RAC, type='l', xlab="Rank", ylab="Abundance") # plot the RAC
```



Looking at the RAC for Site1, we can see that abundance is distributed unevenly among species. In fact, few species have more than 10 individuals and most species have less than 5. This sort of uneven distribution of abundance among species is one of the most ubiquitous patterns in ecology and has provoked a long history of study, theories, and explanations. In fact, the species abundance distribution (SAD) is one of the most widely studied patterns of biodiversity. We will explore another way of illustrating the SAD (other than as a rank-abundance curve) later on in this lesson.

Species Evenness

Now, you may wonder whether all sites in the BCI site-by-species matrix also have highly uneven RACs. But, before moving on to other sites, let's explore aspects of local diversity at Site1 more deeply, that is, by quantifying species diversity in ways other than richness, and by quantifying evenness, dominance, and rarity at Site1.

Species evenness is generally defined as the similarity in abundance among species. Take for instance, the RAC for Site1. While there are many species with similarly low abundances, the few species with relatively high abundances actually drives the pattern to be uneven and the slope of the RAC to be very steep. In fact, one way of quantifying evenness is simply to find the slope of the RAC. This is referred to as

More quantitatively, species evenness reflects the nature of the sample variance. We can see this by considering how the sample variance and one very popular index of evenness (...) are calculated.

Dominance and Rarity

Ken hasn't touched from here on down

Calculate Shannon index. 'Margin = 1' means diversity is calculated row-wise; 'Margin = 2' means diversity is calculated column-wise. With base = exp(1), we are estimating Shannon's index using the natural logarithm of each taxon's relative abundance using the equation; $H' = -\sum(p_i \times \ln(p_i))$


```
shannon <- diversity(BCI, index="shannon", MARGIN = 1, base=exp(1))
```

Can also calculate other diversity metrics

```
simpson <- diversity(BCI, "simpson")
invsimpson <- diversity(BCI, "inv")
fisher <- fisher.alpha(BCI)
```

Let's plot *Pairs is not a normal plotting tool, introduce first*

```
pairs(cbind(shannon, simpson, invsimpson, fisher), pch="+", col="blue")
```



Species richness


```
S <- specnumber(BCI)
```

Rarefaction


```
(raremax <- min(rowSums(BCI)))
```

```
## [1] 340
```

```
Srare <- rarefy(BCI, raremax)
plot(S, Srare, xlab = "Observed No. of Species", ylab = "Rarefied No. of Species")
abline(0, 1)
```



```
rarecurve(BCI, step = 20, sample = raremax, col = "blue", cex = 0.6)
```



Example: <http://www.jennajacobs.org/R/rarefaction.html>

Calculating relative abundances

```
BCI_t <- t(BCI)

BCIrel <- BCI_t
for(i in 1:ncol(BCI_t)){
  BCIrel[,i] = BCI_t[,i] / sum(BCI_t[,i])
}
```

What's one way to test that this worked?

```
colSums(BCIrel)
```

[illegible]